

**INTRODUCTION :**

This Tutorial is based on designing an **Order Processing Database** to deal with Customers placing Orders, firstly at Starbucks and then with Amazon. I have done it this way because Starbucks is quite simple so it provides us with a good starting-point.

Then I look at the questions of extending this simple Database Design to deal with the more complex situation with Amazon. I hope you find this Tutorial interesting and helpful. Please [email me](#) and let me know.

**GETTING STARTED :**

- We stand outside Starbucks, planning to go in.
- We can see Customers, and inside we can anticipate seeing a wide range of things to eat and drink.
- In this Tutorial, I will start by designing a Database for Starbucks and then extend it to include Amazon.
- **My Approach** has three Steps :-
  1. Establish the Scope of the Database
  2. Identify the 'Things of Interest' that are within the Scope, (called **Entities** or Tables in a Database).
  3. Determine the **Relationships** between them.
- At the end of this Tutorial, We have will produced a diagram of the Database Design, which is commonly referred as an 'Entity-Relationship Diagram', or **ERD**
- During the Tutorial, I will refer to Tables as the 'Things of Interest', whereas Data Modellers would refer to Entities, but I think that 'Tables' is more User-Friendly.

**DECIDING THE SCOPE OF OUR DATABASE**

- When we step inside, we see that Starbucks sells a wide range to decide which of them should be included in our Database.
- Right now, we are interested only in something to eat and something to drink.
- Therefore, all the mugs and other items shown in this picture ( **Scope** of our Database, and are not **'Things of Interest'**.

**LOOKING AT THE PRODUCTS :**

- Turning away from the mugs, we can see a display of Food and Drink, and that is what we are looking for.
- Right now, we are thinking about getting something to eat and something to drink.
- When we go in as an ordinary Customer, we are thinking about what we would like to eat and drink.
- However, when we go in as a Database Designer, we look at the display from a different point of view.
- We want to discover the structures in the data and how they are related.



Starbucks Cash Customers (version 1)  
Barry Williams  
1st. August 2008  
DatabaseAnswers.org

Customers\_version1

**COMMENTS :**

- Starbucks (and Retailers in general) know nothing about their

**COMMENTS :**

- Retailers can find out the Name,Address and other details about Customers who pay with a Card.
- Reference data is an important topic and I will come back to it later.
- **Primary Keys**

1. You will notice that the first field in the Customers\_version2 Table is the Customer\_id.
2. It has a 'PK' symbol beside it, which indicates that it is the Primary Key for the Table.
3. The Primary Key is very important and is the way that we can recognise each individual record in the Table.
4. In the case of Starbucks, they will have hundreds of thousands of Customers over a period of time. Therefore we need a way of automatically generating this ID field.
5. The way we do this is to use what is called an '**Auto-Increment**' field. This means that every time you add a new Record, a new Customer ID is automatically generated.
6. Every DBMS has an Auto-Increment field. It is called a 'UniqueIdentifier' data type in Access, an 'Identity' in SQL Server and a 'Sequence' in Oracle.

#### • Foreign Keys

1. When this Primary Key is used in another Table, it is referred to as a 'Foreign Key'.
2. We can see a good example in this diagram, where the customer\_id appears in the Customers\_Payment\_Methods Table as a Foreign Key.
3. This is shown with an 'FK' symbol beside it

#### • Mandatory Key Fields

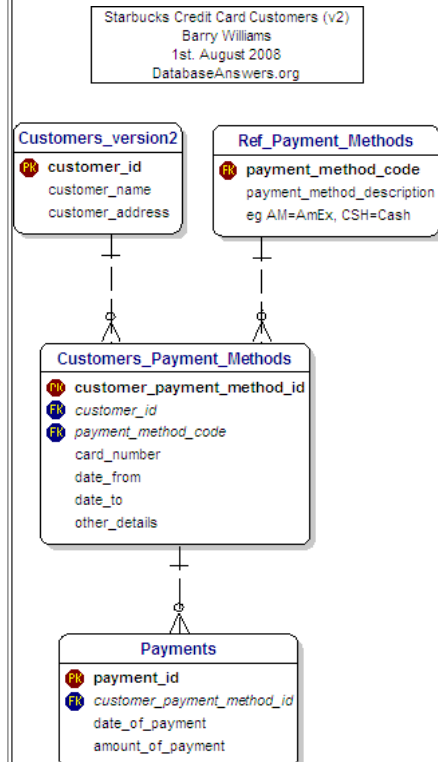
1. A Foreign Key is usually **mandatory**, in other words, a value for a customer\_id in the Customers\_Payment\_Methods Table must correspond to an actual value of the customer\_id in the Customers\_Version\_1 Table.
2. This is shown in the diagram by the short straight line at the end of the dotted line close to the Customers Table.

#### • Optional Key Fields

1. Not every Customer will have a Payment\_Method. In general, they would but we need to allow for situations where Customers change their minds and don't buy anything.
2. In other words, we would say that the Relationship is **optional** at the Customers\_Payment\_Methods Table end.
3. This is shown by the little 'O' at that end of the Relationship dotted line.

#### • One-to-Many Relationships

1. A Customer can have more than one Payment\_Method, for example, American Express or Cash.
2. In other words, we would say that the Relationship is **optional** at the Customers\_Payment\_Methods Table end.
3. This is shown by the symbol that has three small lines at that end of the Relationship dotted line, which is referred to as **Crow's Feet**.



#### COMMENTS :

- We can make a start on the Product hierarchy by looking at the
- At the top level, we have :-
  1. Coffeehouse Favourites
  2. Espresso
  3. Frappuccino

I always like to put things in alphabetical order as soon as possible

Then looking more closely, we can see the Products listed under hand corner, are as follows :-

- Blended Drinks
- Caramel Cream
- Chocolate CremeExpresso
- Strawberries and Cream
- Vanilla Cream


At this point, we have defined three levels within our Product Hierarchy

- The top level 1
  1. Coffeehouse Favourites
  2. Espresso
  3. Frappuccino
    - Blended Drinks
      1. Caramel Cream
      2. Chocolate CremeExpresso
      3. Strawberries and Cream

Then we can see that Caramel Cream, Chocolate CremeExpress Vanilla Cream all have a common Parent Product, which is 'Blend

Products (v1)  
Barry Williams  
3rd. August 2008  
DatabaseAnswers.org

COMMENTS :  
Product Types are in a hierarchy.



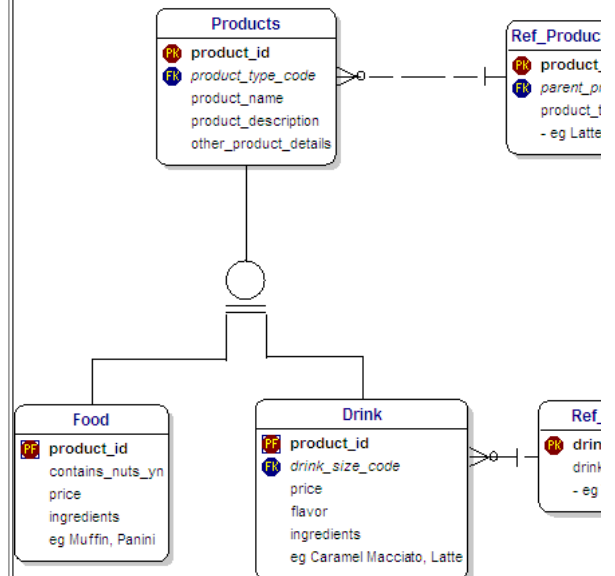
#### COMMENTS :

- This diagram shows how the hierarchies of Products and Product Types discussed are shown in our **Entity-Relationship Diagram**.

#### COMMENTS :

- Food and Drink are specific examples of the more general Thing called a Product.
- They inherit some common attributes from the Product, and also have some of their own.
- For example, Food can contain Nuts but Drink do not contain nuts, but both have a Product Name.
- Rabbits Ears**
- You will notice that the table called 'Product\_Types\_v1' has a dotted line coming out on the right-hand side and going back in again on the top-right corner.
- Data Analysts call this a Reflexive Relationship, or informally, simply 'Rabbits Ears'.
- In plain English, we would say that the Table is joined to itself and it means that a record in this Table can be related to another record in the Table.
- This approach is how we handle the situation where each Product can be in a hierarchy and related to another Product.
- For example, a Product called Panini could be in a Product Sub-Category called 'Miscellaneous Sandwiches' which could be a higher Product Category called 'Cold Food', which itself could be in a higher Product Super-Category called simply 'Food'.
- Next time you go into Starbucks, take a look at the board behind the counter and try to decide how you design the Products area of the Database.
- You should **pay special attention** to the little 'zeros' at each end of the dotted line.
- These are how we implement the fact that the 'Parent Product Type Code' is optional, because the highest level will not have a Parent.
- It's important to think through this level of detail otherwise you will get caught out somewhere down the line when real data can't be stored in your Database.
- In practice, it's smart to have a checklist of things like this to run through before you expose the first draft of your new Database design to a critical audience.
- It's good practice even if you don't have a critical audience because otherwise your design will need to be corrected.
- This topic is discussed again in the final Slide that looks at Deliveries.
- If you find this kind of thing interesting then maybe you would be happy earning a living as a Database specialist.
- If not, then keep well away from it, and make sure your livelihood doesn't depend on your ability to produce a well-designed Database.
- You can trust me when I say that if you don't like it, or have an aptitude for it, it quickly becomes a chore that you would hate, because it requires concentrated thought and attention for prolonged periods of time.
- You could think of it **like solving crosswords** eight hours a day, five days a week.
- Hierarchies in Products and Product Types**
- It's important to understand the difference between these two things, which is that the Product Type Hierarchy defines the Category names and levels, while the Product Hierarchy defines where specific Products fit into the Product Type Hierarchy.

Products (Version 2)  
Barry Williams  
3rd. August 2008  
DatabaseAnswers.org

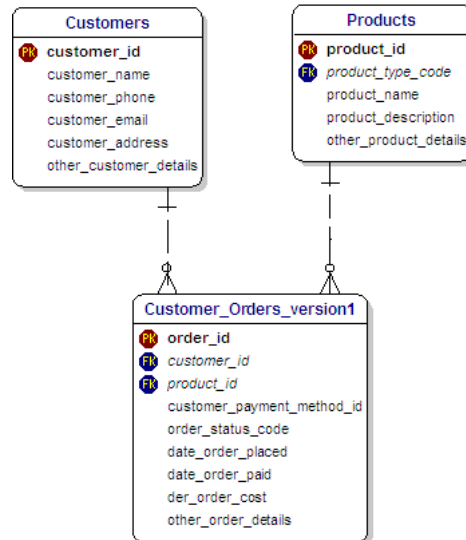


#### COMMENTS :

- This is a very simple (and unrealistic) Diagram.
- It shows that Customers order Products, such as Latte and a Muffin.



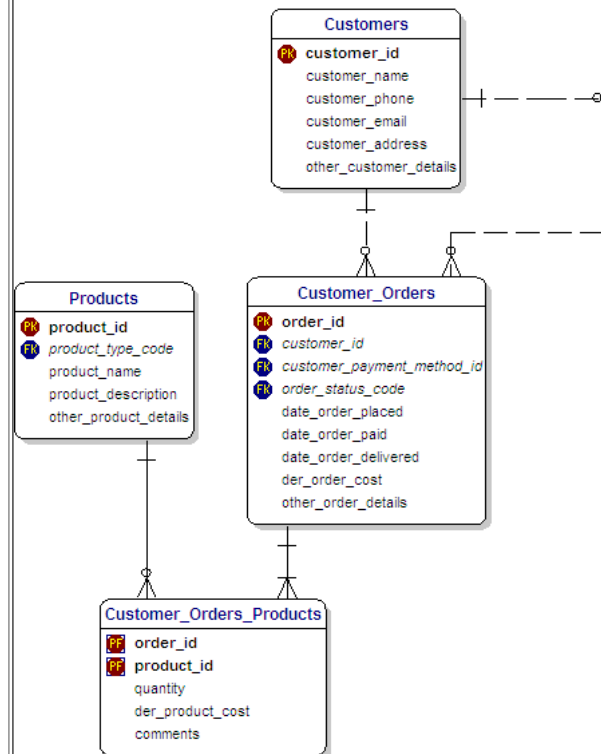
Orders (v1) - Simple (unrealistic) Version  
 Barry Williams  
 3rd. August 2008  
 DatabaseAnswers.org



#### COMMENTS :

- This is a realistic Diagram showing the detail necessary to correctly model Customers, Products and Order details.

Orders (v2)  
 Barry Williams  
 3rd. August 2008  
 DatabaseAnswers.org



#### 1. What have we found in Starbucks ?

- Customers
- Orders
- Products
  - Food
  - Drink

Now we move on to Amazon ...

#### COMMENTS :

- Amazon Customers need to give their Delivery Address, their Home Address. These can all be different.
- We could simply include them all in the Customer record, as w Version 1. However, this looks ugly. It also breaks one of the ve Database Thoery, which is that 'Repeating Groups are not allowe clearly Repeating Groups, and so we must find another way to r

Amazon Customers (version 1)  
Barry Williams  
1st. August 2008  
DatabaseAnswers.org

**Amazon\_Customers\_version1**

- PK** customer\_id
- customer\_name
- customer\_phone
- customer\_email
- billing\_address\_line\_1
- billing\_address\_line\_2
- billing\_address\_city
- billing\_address\_country
- delivery\_address\_line\_1
- delivery\_address\_line\_2
- delivery\_address\_city
- delivery\_address\_state
- delivery\_address\_country
- home\_address\_line\_1
- home\_address\_line\_2
- home\_address\_city
- home\_address\_state
- home\_address\_country
- other\_customer\_details

#### COMMENTS :

- And this is how we do it.
- We have a separate Address Table, which allows us to have more than one Address for any Customer very easily.
- This design also has some other benefits :-
  - We can accommodate more than one person at the same Address. We need to do this because different members of a family may sign-up separately with Amazon.
  - With a separate table of Addresses, we can easily use commercial software to validate our Addresses. To find this kind of software, simply Google for "Address Validation Software". I have used QAS with great success in the past. With this approach, we can always be sure that we have 100% good Address data in our Database.

#### • Reference Data

- This diagram shows Address Types, which are an example of Reference Data.
- This kind of data has the following characteristics :-
  - it doesn't change very much
  - it has a relatively small number of values, usually less than a few dozen and never more than a few hundred.
  - Therefore we can show it with a Code as a Primary Key.
  - Data in Reference Data Tables can be used to populate drop-down lists for Users to select from.
  - In this way, it is used to ensure that all new data is valid.

#### • Standards

- In the Address Table, you will see a field called 'iso\_country\_codes'.
- iso stands for the 'International Standards Organisation'.
- Where possible, it's always good to use national or international standards.


Amazon Customers (version 2)  
Barry Williams  
1st. August 2008  
DatabaseAnswers.org

**Amazon\_Customers\_version2**

- PK** customer\_id
- customer\_name
- customer\_phone
- customer\_email
- other\_customer\_details

**Addresses**

- PK** address\_id
- line\_1\_number\_building
- line\_2\_number\_street
- line\_3\_area\_locality
- city
- zip\_postcode
- state\_province\_county
- iso\_country\_code
- other\_address\_details



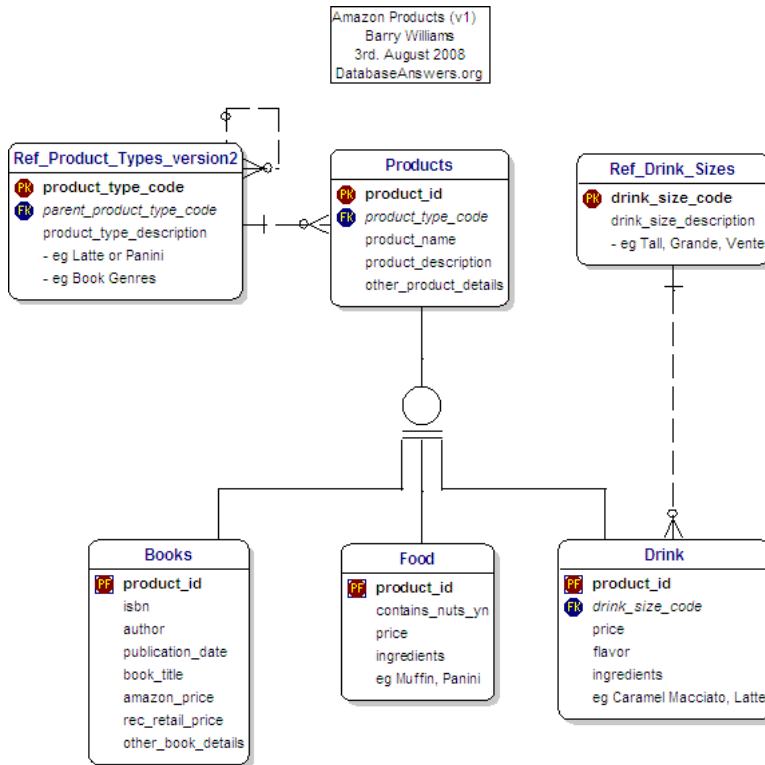
#### COMMENTS :

- Different kinds of Products, like Books, are added quite easily.
- We simply add different kinds of Products, and I have shown B
- You will see that the attributes for Books are specific to Books :
- For example, an Author and an ISBN.
- Because all Products are handled in a Database in a similar way to extend its range of Products.

#### • Inheritance

- The unusual symbol in the middle of the diagram, composed of a circle with a vertical line through it, is how Inheritance is shown using the particular Design, which is called Design.
- Inheritance is a very important topic when you are designing a database.
- In plain English, we would say that Inheritance occurs where a relationship exists between Things of Interest (or Entities).
- You can ask the simple 'Is-a' question - in this case, if we ask 'Is a Book a Product?' clearly the answer is 'Yes' so we think there is an Inheritance relationship.
- In the example of Inheritance shown in this diagram, we can see Books, Food and Drink inheriting from the parent Product.
- However, each type of Product will have specific characteristics.





other types of Products. For example, Books have ISBNs and Authors.

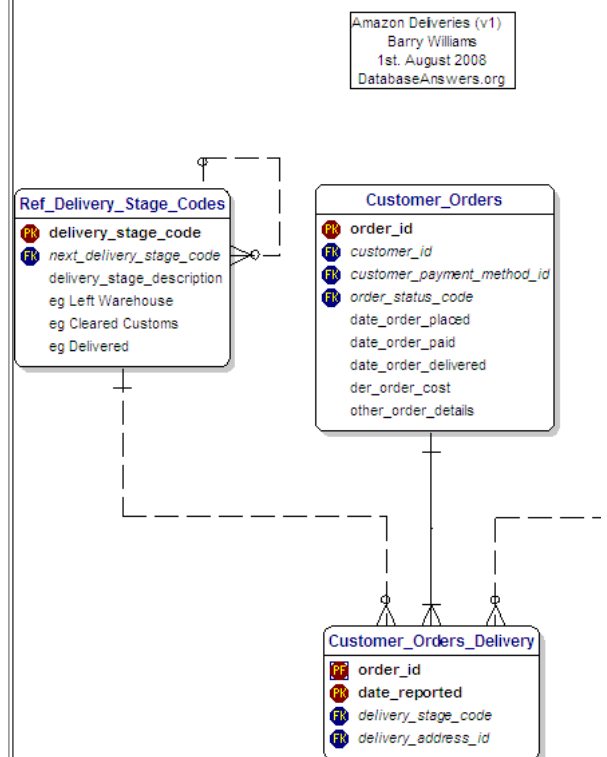
- One of the important things in your Database design is to be consistent with Inheritance relationships.
- However, from many years of experience as a DBA, I should probably not blur in a real physical Database because it can be clumsy to implement.
- I sometimes find myself showing Inheritance in a Logical Data Model when I design the Physical Database, which is what ultimately becomes the Physical Database.

#### • There are three different approaches to implementing Inheritance

1. Implement the design as you see it, in other words, have a separate Product Entity in this diagram.
2. Add all attributes in the Product Super-Type to the Books, Food, and Drink tables. This is appealing because it seems simpler, but more difficulty when it comes to development.
3. Add all attributes from the Books, Food and Drink Sub-Types to the Product table. This results in one large Table, where two-thirds of all Attributes are unused. However, software development can be easier, depending on how it is followed. Currently, wasted space is not such an issue, so this approach is often used.

#### COMMENTS :

- Although Deliveries are complex in the real world, it is easy to add them to our Database design.
- We simply add a Reference Data Table of Delivery Status Codes and one Table for tracking Order Deliveries.
- You will notice that the name of this Table is 'Customer\_Orders\_Delivery', which follows our convention of combining the names of Tables that are combined in the data that is stored in the Table.
- **Rabbits Ears**
- You will notice that the table called 'Ref\_Delivery\_Stage\_Codes' has a dotted line coming out on the right-hand side and going back in again on the top-right corner.
- Data Analysts call this a Reflexive Relationship, or informally, simply 'Rabbits Ears'.
- In plain English, we would say that the Table is joined to itself and it means that a record in this Table can be related to another record in the Table.
- This approach is how we handle the situation where each Delivery Stage is related to another one.
- When you buy a product from Amazon, you can track its progress as it leaves the Warehouse, clears Customs and is finally delivered to you.
- You should **pay special attention** to the little 'zeros' at each end of the dotted line.
- These are how we implement the fact that the 'Next Delivery Stage Code' is optional, because the last Stage in Delivery will not have a Next one.
- It's important to think through this level of detail otherwise you will get caught out somewhere down the line when real data can't be stored in your Database.
- In practice, it's smart to have a checklist of things like this to run through before you expose the first draft of your new Database design to a critical audience.
- It's good practice even if you don't have a critical audience because otherwise your design will need to be corrected.



#### Finally, here's a Summary of the Rules for designing a Database

1. Define the **Scope** of your Database.
2. Define the **"Things of Interest"**, (e.g. **Customers** and **Order**), that are within Scope.
3. Establish how these Things are related and write down the 'Business Rules'.  
For example, "A **Customer** can have zero, one or many **Addresses**."
4. Determine what else you know about these Things.  
For example, "Books have an ISBN and one or many **Authors**."
5. Identify the Reference Data, such as **Address Types** and **Product Types**.

- You need to define a Primary Key for all Tables.
  - For Reference Tables, use the 'Code' as the Key, often with only one other field, which is the Description field.
  - For all other Data, you can use a generated number as the Primary Key.
- This has some major benefits, for example, it provides flexibility, and it's really the only choice for a Database supporting a Web Site.

#### • How to Validate your Database Design

It's always useful to validate your design.

One good way is to talk through it with someone who knows a little about what you are doing, or about Amazon or Starbucks !!

Another good way is to check that you will be able to store actual data in your new Database.

- Firstly, obtain a small set of Sample Data by thinking about your last trip to Starbucks, or the last book you bought from Amazon.
- Confirm the first draft of the Database design against the Sample Data.
- Obtain from the users some representative enquiries for the Database, e.g. "How many Hot Drinks are on offer in Starbucks ?"
- Review Code or Type Data which is (more or less) constant, which can be classified as **Reference Data**. Look for external standards which can be national or international. For example, Currency or Country Codes might have ISO Codes.

• Finally, define User Scenarios and step through them with some sample data to check that that Database supports the required functionality.

• **Please email me with your comments**

• I hope you have found this Tutorial interesting and useful.

• Please [email me](#) with your questions or suggestions so I can improve this first draft Tutorial.

Good luck with designing your first Database !

**Barry Williams**

Principal Consultant  
Database Answers Ltd.  
London, England

[Home](#)[Ask a Question](#)[Best Practice](#)[Communities](#)[Contact Us](#)[Data Models](#)[Search](#)[Site Map](#)

© DatabaseAnswers.org 2008